

Online Appendix

Semiparametric Value-at-Risk and Expected Shortfall with a Real-Time Misspecification
Score

Sean Qin

September 2026

OA.1 Glossary of acronyms

ARCD	autoregressive conditional density model (Hansen, 1994)
CAViaR	conditional autoregressive value-at-risk (Engle and Manganelli, 2004); SAV, its symmetric absolute value form
CC, UC	conditional coverage (Christoffersen, 1998) and unconditional coverage (Kupiec, 1995) exception tests
CRPS	continuous ranked probability score
CRSP, WRDS	Center for Research in Security Prices; Wharton Research Data Services
DM	Diebold–Mariano test (Diebold and Mariano, 1995)
EM, IG, OAS	emerging markets; investment grade; option-adjusted spread
ERM	empirical risk minimization
ES, VaR	Expected Shortfall; Value-at-Risk
EVT, GPD	extreme-value theory; generalized Pareto distribution
EWMA	exponentially weighted moving average variance (J.P. Morgan/Reuters, 1996)
FHS, HS	filtered historical simulation (Barone-Adesi et al., 1999); historical simulation
FRTB	Fundamental Review of the Trading Book (Basel Committee on Banking Supervision, 2019)
FWER	family-wise error rate
FZ0	the zero-homogeneous Fissler–Ziegel joint (VaR, ES) scoring function (Fissler and Ziegel, 2016; Patton et al., 2019)
GARCH, GJR	generalized autoregressive conditional heteroskedasticity (Bollerslev, 1986); the Glosten–Jagannathan–Runkle leverage form (Glosten et al., 1993)
GAS	generalized autoregressive score model (Patton et al., 2019)
GBC	generative Bayesian computation (Polson and Sokolov, 2023)
GBM	gradient boosting machine (gradient-boosted regression trees)
GFC	the 2008 global financial crisis
IQN	implicit quantile network (Dabney et al., 2018)
MCS	Model Confidence Set (Hansen et al., 2011)
PIT	probability integral transform
PZC	Patton, Ziegel and Chen (Patton et al., 2019)
QML	quasi-maximum likelihood
RV	realized variance
SPA	superior predictive ability test (Hansen, 2005)
TAQ	NYSE Trade and Quote database

OA.2 Algorithms

Algorithm OA.1 (amortized shape model). (i) Assemble the pooled panel $\{(s_{j,t}, \hat{z}_{j,t})\}$ over names j and days t . (ii) Train H_ϕ by SGD on $\mathbb{E}_{\tau \sim \lambda} \rho_\tau(\hat{z} - H_\phi(s, \tau))$ [network], or boost per- τ regressions on a grid [trees]; enforce monotonicity by rearrangement. (iii) For any asset, including one with no return history, using characteristics-only s , evaluate $H_\phi(s_t, \cdot)$; no per-asset estimation step exists. (iv) To *sample*, draw $\tau^{(m)} \sim U(0, 1)$ and return $z^{(m)} = H_\phi(s_t, \tau^{(m)})$.

Algorithm OA.2 (the misspecification meter, as deployed). (i) Fit equation (2) of the paper on a trailing window; store $\hat{z}_{t-62}, \dots, \hat{z}_t$. (ii) Compute mk_{63} , $|\text{sk}_{63}|$, J_5 by equation (9) of the paper. (iii) Convert each to its percentile against the trailing pooled panel (same asset class); the score is the maximum of the three percentiles, re-ranked into deciles. (iv) Act by the paper: top decile $\Rightarrow$ the parametric model’s average pinball loss runs $\sim 2\text{--}3\%$ higher within this score region (a conditional average, not a per-name point estimate), flag the asset for the nonparametric / residual-hybrid engine and for model-risk review; elsewhere $\Rightarrow$ no action (in the amortized setting the engine is run regardless; the meter serves surveillance, and day-level switching is unsupported, Section 4.3 of the paper). (v) Latency: detection lags episode onset by median 0 / mean 2–4 days; the edge is back-loaded within episodes, so the lag forfeits little.

Deployed pipeline. Each day: (i) update each name’s GARCH filter state $\hat{\sigma}_{t+1}$ (recursion only; parameters refit annually); (ii) assemble the state vector s_t and evaluate the shared shape model $q_\phi(\tau \mid s_t)$, one model for the whole book, no per-name estimation; (iii) apply the frozen GPD tail below p_0 and, where the desk elects the coverage overlay, the adaptive conformal shift c_τ , updated from the previous day’s panel breach frequency; (iv) report $\hat{Q}_{t+1}(\tau) = \hat{\mu} + \hat{\sigma}_{t+1}(Q_{t+1}^z(\tau) + c_\tau)$ at the desk’s levels, with VaR and ES_{97.5} both read from the coherent min-envelope curve Q^z of equation (5) of the paper (VaR at the α node, ES its numerical tail integral); (v) read the misspecification score equation (9) of the paper as the model-risk monitor: top decile means the parametric fallback’s conditional average loss in that score region runs $\sim 2\text{--}3\%$ higher. *Annually:* refit filters, retrain the shape model, re-estimate the tail, all on trailing data. *New listing:* run the separate characteristics-only amortized forecaster until an own-history sample sufficient to initialize the per-name scale filter accrues; then attach this residual-hybrid pipeline. Total marginal cost per name per day: one filter update and one model evaluation.

OA.3 Supplementary figures




Figure OA.1: Timing diagnostics for the frontier. Top-decile edge under four alignments of the mk_{63} score with the daily loss (200 CRSP names; per-date DM 5.5–6.5 in all four). The deployed lag dominates, and the contaminated same-day alignment yields less, the opposite of what a look-ahead artifact would produce.




Figure OA.2: Exception tests restated per asset (200 CRSP names, 99% VaR, the rows of Table 4 of the paper): share of names whose individual Kupiec test is consistent with nominal. HS over-covers (conservative); EWMA fails; the EVT-tailed engine matches the parametric benchmarks while beating them on loss and the joint FZ0 score (Table 4 of the paper). Date-clustered tests in the table notes.




Figure OA.3: The edge grows with horizon. Direct-model/GARCH pinball ratio at $h = 1, 5, 10, 20$ days (150-name design panel; training labels whose h -day window crosses the split are purged): the $\sqrt{h}$ rule holds the one-day variance flat across the window while the fitted GARCH variance mean-reverts, and the gap grows with the horizon, from 0.45% at $h=1$ to 2.1% at $h=20$; at $h=10$ an iterated-variance GARCH- t benchmark removes about half of it (main text, ten-day subsection). The forecaster is the direct multi-day model of the main text’s ten-day subsection, *not* the residual-hybrid engine, whose architecture is one-day.




Figure OA.4: The FRTB battery. Left: average pinball loss by model (axis truncated to resolve the ordering; differences carry DM 5.6–9.6 vs the hybrid, all $p \approx 0$); CAViaR, not shown, statistically ties the hybrid. Right: ES_{97.5} predicted (the tail integral of Table 4 of the paper) vs realized below each model’s own VaR, a model-dependent diagnostic, not a ranking (every fat-tailed entrant’s integral exceeds its own-tail mean); the joint FZ0 score of the paper is the ranking criterion.




Figure OA.5: The frontier across universes: nonparametric edge vs post-GARCH residual kurtosis, tier means and single assets with study-reported kurtosis only (country universe corr 0.53; cross-asset corr 0.43). The hyperinflation-FX tier (kurtosis ~2000–3300) carries the only decisive standalone wins; high kurtosis is predictive but neither necessary nor sufficient: carbon (288) and corn (20) tie, and NGN’s win is not significant. The Korea-2026 control (kurtosis 4–13, GARCH wins all 23 assets) sits in the lower-left regime.




Figure OA.6: The misspecification frontier. Nonparametric pinball edge over GARCH- t by decile of the trailing 63-day residual excess kurtosis (left) and residual asymmetry (right), 200 CRSP names, 221.6k test asset-days. Deciles 1–8 lie within $\pm 0.2\%$ (values from `composite_profiles_results.json`); the edge jumps in the top decile (+2.98% kurtosis, +2.78% asymmetry; per-date DM 10.5 in the top kurtosis decile).

OA.4 Supplementary tables

Direct expected-shortfall backtests

Exception counts test the VaR level; they do not test ES. The main text adds two direct ES *calibration* backtests: the McNeil–Frey (2000) exceedance-residual test, whose statistic is the mean standardized loss beyond VaR, zero under a correct ES (McNeil and Frey, 2000); and the Acerbi–Székely (2014) statistic Z_2 , negative when ES understates the realized tail (Acerbi and Székely, 2014). These complement the jointly consistent FZ0 loss used for forecast comparison; Nolde and Ziegel (2017) and Bayer and Dimitriadis (2022) draw the distinction between backtesting the calibration of an ES forecast and comparing forecasts through a jointly consistent score. Table OA.5 reports the two calibration tests on the same 200-name panel and 221,600 held-out asset-days (1,108 test dates) as the full-panel FZ0 re-scoring, with the engine’s VaR and ES read from its own coherent min-envelope curve Q^z of equation (5) of the paper, the forecasts the joint loss scores, and inference by a stationary block bootstrap over calendar dates. Gaussian innovations are decisively inadequate in this daily-data battery: the Gaussian filter is worst on every ES metric, its mean exceedance residual reaching -0.77 standard deviations at the 1% level ($Z_2 = -0.99$). Among the fat-tailed entrants no model dominates every diagnostic. On the Acerbi–Székely Z_2 the engine is best calibrated in the tail: it carries the statistic closest to zero at 1% (-0.05 , not rejected) and is tied with GARCH- t at 2.5% (both -0.10). The engine and GARCH- t are the two entrants that pass the 1% Kupiec unconditional-coverage test; FHS and the Gaussian filter fail it. On the McNeil–Frey exceedance residual the engine sits closest to zero at 1% and is the clearest non-rejection (-0.06 , $p = 0.38$), ahead of FHS, which fits the empirical residual tail directly (-0.10 , $p = 0.07$), while GARCH- t and the Gaussian filter are rejected; at 2.5% FHS, GARCH- t and the engine all fall inside the non-rejection region. The reading is consistent with the exception tests and the joint score: fat tails materially improve tail calibration, the engine and GARCH- t are the best-calibrated pair on the direct tests, and the engine’s decisive margin is the comparative FZ0 loss, where it beats both GARCH- t and FHS at each level.

Realized measures and the scale–shape decomposition

The main-text battery reads only daily returns, the regime the estimator is built for. To place it against a method that reads intraday data, and to ask what the misspecification frontier is made of, we build five-minute realized measures from TAQ for thirty large-capitalization U.S. equities over 2014–2024 (realized variance is the sum of squared five-minute log-returns, cleaned by regular sale conditions, a daily price band, and a 20% return clip) and fit a Realized GARCH (Hansen et al., 2012): the log-linear conditional variance $\log h_t = \omega_R + \beta_R \log h_{t-1} + \gamma_R \log \text{RV}_{t-1}$ with a measurement equation $\log \text{RV}_t = \xi_R + \phi_R \log h_t + \tau_R(z_t) + u_{R,t}$ linking the realized measure to the

Table OA.1: Summary of the paper’s results and scope conditions. Edges are percentage pinball reductions unless noted; DM statistics are per-date (conventions note of Section 3 of the paper), one-sided. Panel A: the deployed amortized residual-hybrid engine (worst case on the accuracy yardstick: no detectable loss; calm-cell 90% lower confidence bounds sit above -0.25%). A separate block reports the direct multi-day extension, which is not the engine. Panel B: boundary evidence from configurations not deployed, which motivates the design.

<i>Panel A: the engine</i>		
Result	Magnitude	Significance
vs GARCH- t / EWMA / HS, $h=1$ (avg)	+0.3 / +1.5 / +1.9%	DM 4.4–8.3
Top misspecification decile	+2.98%	DM 10.5
Hyperinflation FX (ARS, EGP)	+12–14%	DM 2.9–3.8
Frontier equity indices (4 of 8)	$\sim +1$ –3%	DM 2.1–3.5
2000–2013 frozen-spec holdout (untouched era)	top decile +0.9%	DM 2.2 (sort); recursive real-time rule +1.0%, DM 2.15; frozen-constant restatement in text
Point-in-time universe (Q1-2000 top 200, delistings kept)	top bucket +2.9%; overall +0.6%	DM 6.6 / 9.0, frozen real-time thresholds
Strict calendar splits; annual-refit walk-forward	top-decile edge holds under all three	DM 6.2 (2020 cut); 3.1 (overall, pre-GFC cut); 4.41 (annual-refit residual-hybrid walk-forward, +2.12%)
Familywise error over the 3×10 signal-decile grid	9 cells survive 5% FWER	Romano–Wolf step-down, incl. all three top deciles
Exception battery through the 2008 crisis	Kupiec 94% / Christof. 95% at 99%	date-clustered t 0.17 at 99%; at 97.5%, 0.35 (EVT tail) and -0.27 to -0.44 (adaptive overlay); HS, EWMA fail
Joint (VaR, ES) FZ0 score, full panel	accuracy layer lowest at 1%, never significantly beaten at 2.5%	DM 4.9 vs GARCH- t and FHS at 1%, 5.5 vs GARCH- t at 2.5%; 8.8–9.9 vs PZC GAS, 3.8-to-tie vs Taylor ES-CAViaR (multi-start); tie group with GJR-skew- t and Taylor at 2.5%; adaptive overlay within noise of the accuracy layer (top-decile cost at 2.5%)
Cold start / young listings	+9% at 15–30 days; +6.5–7% to one year	vs an expanding-window Gaussian at levels 0.05–0.95; point-in-time characteristics; first-week VaR not calibrated
Calm-quintile edge	≈ 0 (point estimates +0.0–+0.2%)	bottom-decile 90% CIs lie above -0.25% ; no detectable loss
SAV-CAViaR (shares the frontier)	tie on its common sample (both in 90% MCS); within a third of a percent on the frontier rows	DM 0.4 (Table 4 of the paper); -1.8 top decile, -1.6 overall, engine over CAViaR (Table OA.10)
<i>Extension (a separate direct multi-day model, not the engine)</i>		
Ten-day horizon vs $\sqrt{h}$ -scaling	+1.4% (2014–24); +0.5–0.8% vs iterated GARCH- t	holds in 2020–22 stress; about half is the $\sqrt{h}$ rule (Table OA.13); 2000–13 rerun within noise, DM at most 0.3
<i>Panel B: boundary evidence from configurations not deployed</i>		
Bare standalone neural net, single names	-4.5 to -6.2%	DM +12 to +18; motivates the hybrid design
Standalone nonparametric, calm single assets	GARCH as good or better	pooled DM -9.5 to -13.7 ; the shape edge needs misspecification
Korea 2026 crisis, 23 assets (falsification case)	GARCH wins all	DM -2.5 to -5 ; a price crash is not a shape break
mk ₆₃ -reweighted FHS (reweighting probe)	worse than plain FHS	DM -11 overall and top decile; sharp univariate conditioning collapses the sample
Baltic freight; IG credit ^a	+15%; +52%	DM 9.6; 30.6; accuracy only, calibration open

^a Credit and freight fail breach-independence (exceptions cluster); their entries are accuracy results only. The credit magnitude is from the FRTB-grade single-asset battery; the smaller 4% figure in the paper’s cross-universe table is from the separate lean hybrid-vs-GARCH sweep, different battery, both reported.

Table OA.2: Double sort on the misspecification score and trailing volatility. Cells are the nonparametric-over-GARCH- t pinball edge (%) in independent quintiles of the misspecification score (rows) and trailing volatility (columns), design-era test panel.

	rv ₆₃ Q1 (calm)	Q2	Q3	Q4	Q5 (active)
mk ₆₃ Q1 (low)	+0.24	+0.34	+0.10	+0.01	−0.14
Q2	+0.16	+0.20	+0.22	+0.09	−0.57
Q3	−0.02	+0.19	+0.25	+0.44	+0.04
Q4	+0.08	+0.15	+0.35	+0.23	+0.10
Q5 (high)	−0.10	+0.44	+0.68	+0.61	+2.33

Top-row date-clustered DM by volatility quintile: 0.7, 2.3, 1.2, 2.8, 6.4.
The bottom-left and top-right regions are the point: neither ingredient alone produces the concentration.

Table OA.3: Model selection by task and regime. Each row is established in the section cited.

Task / regime	Winner	Where
Point quantile accuracy, tabular state	GBM (trees)	Section 2.3 of the paper
Single asset, calm (~90% of asset-days)	GARCH- t	Section 4 of the paper
Single asset, top misspecification decile	Nonparametric (+2.98%, DM 10.5)	Section 4 of the paper
Amortized cross-name panel, any regime	Nonparametric always; no gate	Section 4.3 of the paper
Regulatory ES with a $\geq$ GARCH floor	Residual-hybrid + EVT + adaptive overlay	Section 5 of the paper

latent variance, estimated jointly by Gaussian quasi-maximum likelihood. Table OA.6 reports the result on 1,108 common test dates. Two facts follow. The realized scale lowers FZ0 well below the daily-data accuracy core at both levels (DM 6.2 at 2.5%, 4.5 at 1%): intraday information prices volatility more sharply than any daily filter, and the economic gap is largest in the top misspecification decile, where the daily scale is coarsest. But the gain is scale, not shape. Holding the realized scale fixed and re-fitting the pooled state-conditioned residual-quantile learner, the EVT tail, and the misspecification score on the realized residuals moves FZ0 within noise: every stage’s Diebold–Mariano statistic against the realized benchmark is at most 1.4 in magnitude, overall and in the top score decile, the largest being the shape learner’s −1.4 at 1% (directional, not significant). On these large-capitalization names a realized scale absorbs most of what the score detects on daily-GARCH residuals, so on the intraday-covered segment the two are substitutes; the score’s distinctive value is in the daily-data regime and the cross-asset and day-one settings where a realized scale does not exist.

Table OA.4: Cross-universe multiplicity: Holm step-down over all 93 instrument-level comparisons in the three exploratory universes of the paper’s cross-universe table (24 FX pairs, 26 country indices, 43 cross-asset instruments). One-sided p from each instrument’s date-clustered DM statistic; familywise $\alpha = 5\%$. Ranks below twelve all have $p > 0.02$ and are omitted.

Rank	Instrument	Universe	DM	p (one-sided)	Holm cutoff	Verdict
1	PJM power ^a	cross-asset	283.1	$< 10^{-16}$	5.38×10^{-4}	artifact ^a
2	Baltic Dry	cross-asset	9.36	$< 10^{-16}$	5.43×10^{-4}	survives
3	VIX	cross-asset	5.71	5.7×10^{-9}	5.49×10^{-4}	survives
4	V2X	cross-asset	5.03	2.5×10^{-7}	5.56×10^{-4}	survives
5	VVIX	cross-asset	4.66	1.6×10^{-6}	5.62×10^{-4}	survives
6	USDEGP	FX	3.80	7.2×10^{-5}	5.68×10^{-4}	survives
7	Kenya NSE20	indices	3.52	2.2×10^{-4}	5.75×10^{-4}	survives
8	Wheat	cross-asset	3.51	2.2×10^{-4}	5.81×10^{-4}	survives
9	MOVE	cross-asset	3.22	6.4×10^{-4}	5.88×10^{-4}	fails
10	Sri Lanka CSEALL	indices	3.20	6.9×10^{-4}	5.95×10^{-4}	fails
11	Gasoline	cross-asset	3.20	6.9×10^{-4}	6.02×10^{-4}	fails
12	USDARS	FX	2.86	2.1×10^{-3}	6.10×10^{-4}	fails

The family is every instrument-level look in the three exploratory universes; the Korea-2026 falsification case (23 assets, GARCH- t wins throughout) is not part of the win family. Bonferroni at 0.05/93 gives the identical survivor set. The pooled crisis-FX comparison (DM 4.12, $p \approx 1.9 \times 10^{-5}$) and the kurtosis-edge correlations (0.53 across indices, 0.43 cross-asset) are pooled and gradient statistics outside this instrument-level family; they are the replication evidence the paper’s cross-universe claims rest on.

^a Retained in the family because the raw sweep contains it, but not a win: the paper’s cross-asset audit reclassifies electricity as a log-return artifact on near-zero and negative prices, refit on price differences, GARCH wins PJM (DM -2.73), so the survivor count the text quotes is seven.

Table OA.5: Direct expected-shortfall backtests. 200 CRSP names, pooled test panel, the same held-out rows as the full-panel FZ0 re-scoring; the engine’s VaR and ES are its own coherent min-envelope forecasts. Breach is the realized exception rate (target α); Kupiec p is the pooled unconditional-coverage test. Z_2 is the Acerbi–Székely (2014) statistic, negative when ES understates the tail. MF is the McNeil–Frey (2000) mean exceedance residual (standardized loss beyond VaR, zero under correct ES) with its bootstrap p .

Model	$\alpha = 2.5\%$				$\alpha = 1\%$			
	Breach	Kupiec p	Z_2	MF (p)	Breach	Kupiec p	Z_2	MF (p)
GARCH-normal	0.027	0.00	−0.31	−0.55 (0.00)	0.016	0.00	−0.99	−0.77 (0.00)
GARCH- t	0.027	0.00	−0.10	−0.05 (0.14)	0.010	0.17	−0.09	−0.23 (0.00)
GARCH-FHS	0.031	0.00	−0.24	+0.01 (0.84)	0.012	0.00	−0.29	−0.10 (0.07)
Engine (EVT-tailed)	0.028	0.00	−0.10	+0.02 (0.58)	0.010	0.07	−0.05	−0.06 (0.38)

Inference for the McNeil–Frey and Z_2 p -values is a stationary block bootstrap over the 1,108 trading dates (mean block ten dates) rather than treating the 221,600 asset-days as independent, since the 200 names share within-day market shocks. Kupiec is the pooled analytic unconditional-coverage test, hypersensitive at this sample size (every model is rejected at 2.5%). The engine’s ES is the converged integral of Section 2.2 of the paper. The static conformal shift of earlier versions, estimated on the body and applied to the engine, over-covers (breach 1.66%, $Z_2 = +0.31$, McNeil–Frey $p = 0.002$ at 2.5%); the adaptive overlay that replaces it is reported in Section 5.1 of the paper. Source: `engine_esbt_results.json`.

Table OA.6: The scale–shape decomposition. Thirty large-cap U.S. equities, 2014–2024 (1,108 test dates), joint (VaR, ES) FZ0 loss (lower is better) at both FRTB levels. The scale is a proper Realized GARCH except in the first row (the daily-data accuracy core). DM is date-clustered (Newey–West, 5 lags) against the realized scale with an unconditional residual quantile; a positive DM is worse than that benchmark.

Model	FZ0 2.5%	DM	FZ0 1%	DM
Daily core (GARCH- t + EVT)	1.500	+6.2	1.774	+4.6
Realized scale + residual quantile	1.445	—	1.715	—
+ EVT tail	1.445	−0.1	1.716	+1.3
+ pooled shape learner	1.442	−0.8	1.708	−1.4
+ shape + EVT (realized-hybrid)	1.447	+0.7	1.717	+0.5

All five rows are computed in one pipeline on the same test rows. ES is exact wherever a closed form exists: the two GPD-tailed rows use the closed-form tail integral and the unconditional row the empirical tail mean. The shape row is a body-only forecast with no tail model below its lowest fitted level, so its ES keeps a flat extension there and is a lower bound on severity; the realized-hybrid row uses the converged integral of its min-envelope curve. The daily core’s DM is against the realized benchmark (the intraday-scale gain). The pooled shape learner is the paper’s state-conditioned residual-quantile model, retrained on the realized-scale residuals. Returns are CRSP daily; the intersection is common across models. Source: `scaleshape_canonical_results.json`.

OA.5 Standalone GARCH-EVT on the canonical rows

The manuscript’s comparison set omits standalone GARCH-EVT (McNeil and Frey, 2000), which Kuester et al. (2006) rank first among classical methods. This section reports it on exactly the rows, filter and splits of the frontier sorts (Table 1 of the paper) and of the full-panel FZ0 run (Figure 2 of the paper), from `job_garch_evt.py` (result file `garch_evt_results.json`). Two forms are run: per name, with the GPD tails fitted to each name’s own training residuals, and pooled, with one threshold and one (ξ, β) per tail across names. The pooled body row is the frontier engine of Table 1 (no EVT branch); the engine row adds the body/EVT minimum and rearrangement of equation (5). Table OA.7 gives the eleven-level pinball comparison and Table OA.8 the joint score.

Table OA.7: GARCH-EVT on the frontier rows: pinball edge by score region and head-to-head comparisons.

Model	Top mk ₆₃ decile	Deciles 1–9	Top composite decile	Overall
<i>Edge over GARCH-t, % (per-date DM)</i>				
GARCH-EVT, per name	+0.31 (3.62)	−0.22 (−6.61)	+0.37 (5.35)	−0.16 (−5.00)
GARCH-EVT, pooled tail	−0.16 (−2.69)	−0.12 (−4.55)	−0.06 (−1.40)	−0.12 (−4.89)
GARCH + pooled body (no EVT)	+2.98 (10.48)	+0.02 (0.43)	+2.79 (8.91)	+0.35 (5.35)
Engine (body/EVT minimum)	+2.48 (10.08)	−0.01 (−0.25)	+2.19 (8.00)	+0.27 (4.62)
<i>Head-to-head, % (per-date DM)</i>				
Engine over GARCH-EVT pooled	+2.64 (8.91)	+0.10 (2.47)	+2.24 (7.51)	+0.39 (7.66)
Engine over GARCH-EVT per name	+2.18 (9.63)	+0.21 (4.63)	+1.82 (7.39)	+0.43 (8.22)
Pooled body over GARCH-EVT pooled	+3.14 (9.39)	+0.14 (2.84)	+2.84 (8.37)	+0.47 (7.96)
Engine over pooled body	−0.51 (−9.04)	−0.04 (−2.28)	−0.62 (−9.64)	−0.09 (−5.25)

200 names, 221,600 test rows, same rows and splits as Table 1 of the paper and Figure 2 of the paper. Edge = (reference pinball − model pinball)/reference pinball over the eleven-level grid; DM is the per-date Newey–West(10) Diebold–Mariano statistic. GARCH-EVT is McNeil–Frey: GARCH(1,1)- t filter, GPD tails fitted by maximum likelihood to the 10% most extreme training residuals per tail (per name, or pooled across names), empirical residual quantiles between the thresholds. The engine’s pooled GPD at $p_0 = 0.025$ has $\hat{\xi} = 0.298$, $\hat{\beta} = 0.639$ on 6075 exceedances; the EVT branch is the tighter of the two at the 1% node on 60.1% of test rows and at the 2.5% node on 59.5%.

Threshold diagnostics. Over a threshold grid $p_0 \in \{0.025, 0.05, 0.10, 0.15\}$ the pooled lower-tail fit gives $\hat{\xi} = 0.298, 0.269, 0.204, 0.164$ on 6,075, 12,150, 24,300, 36,450 exceedances; per name at $p_0 = 0.10$ the interquartile range of $\hat{\xi}$ is 0.058 to 0.281 (median 0.174) with a median of 125 exceedances, and no name has $\hat{\xi} \geq 1$. The EVT branch is the tighter of the two inside the engine’s minimum on 60.1% of test rows at the 1% node and 59.5% at the 2.5% node. On the 2000–2013 holdout rows (`job_holdout_garch_evt.py`) the frozen learner beats the pooled GARCH-EVT by +1.16% (DM 2.52) in the top holdout decile and by +1.30% (DM 1.38) in the frozen design-era bucket, and by −0.01% (DM −0.31) in deciles one through nine.

Additional evaluation statistics on the same rows

The same run adds four statistics that go beyond mean loss. (i) The conditional-predictive-ability test of Giacomini and White (2006), with the asset-day loss differential $d_{it} = L_{it}^{\text{GARCH-}t} - L_{it}^{\text{engine}}$

Table OA.8: GARCH-EVT on the full-panel FZ0 rows.

Model	$\alpha = 1\%$			$\alpha = 2.5\%$		
	FZ0	Breach	DM vs engine	FZ0	Breach	DM vs engine
Engine (body/EVT minimum)	2.1465	1.04%	—	1.8486	2.78%	—
Engine + static shift (97.5%)	2.1465	1.04%	0.00	1.8664	1.66%	2.57
GARCH + pooled body (no EVT)	2.1411	1.32%	-1.09	1.8432	3.25%	-1.54
GARCH-EVT, pooled tail	2.1551	1.06%	4.51	1.8583	2.89%	6.33
GARCH-EVT, per name	2.1693	1.13%	5.14	1.8596	2.93%	5.36
GARCH(1,1)- t	2.1557	1.03%	4.94	1.8542	2.71%	5.50

FZ0 is the zero-homogeneous Fissler–Ziegel joint (VaR, ES) loss, lower is better; DM > 0 means the row model scores worse than the engine (per-date Newey–West(10)). The engine row is the accuracy layer (no conformal shift) at both levels and is the reference of every DM; the static-shift row adds the split-conformal shift of earlier versions at 97.5%, estimated on the body’s calibration errors, and coincides with the engine at 1%; Stage 4 now uses the adaptive overlay of Section 5.1 of the paper. GARCH-EVT ES is the McNeil–Frey closed form; the engine ES is the converged integral of the monotonized quantile curve (Section 2.2 of the paper), and the body ES a 20-node integral, since the body has no quantile below its lowest fitted level. Top-mk₆₃-decile DM of pooled GARCH-EVT versus the engine: 2.26 at 1% and 3.42 at 2.5%.

(eleven-level pinball) and instruments $(1, \text{score}_{i,t-1})$; the Wald statistic uses a Newey–West(10) covariance of the per-date sums. (ii) Murphy diagrams (Ehm et al., 2016): the mean elementary quantile score $S_\theta(q, y) = (\mathbb{I}\{y < q\} - \alpha)(\mathbb{I}\{\theta < q\} - \mathbb{I}\{\theta < y\})$ of the accuracy layer minus that of a competitor, on a grid of forty return-space thresholds θ (the 0.1% to 15% quantiles of test returns); a negative value means the accuracy layer is better at that θ , and a curve that is negative everywhere would mean dominance under every consistent scoring function for the α -quantile. (iii) The dynamic quantile test of Engle and Manganelli (2004) per name with four hit lags and the VaR. (iv) The dispersion of the per-date loss differential.




Figure OA.7: Murphy diagrams for the VaR forecasts on the frontier rows: mean elementary quantile score of the accuracy layer minus that of GARCH- t (solid) and of the pooled GARCH-EVT (dashed), against the threshold θ . Negative values favor the accuracy layer. It is better on 70% (GARCH- t) and 85% (GARCH-EVT) of the grid at $\alpha = 1\%$ and on 65% and 85% at $\alpha = 2.5\%$; the competitors win in a band of thresholds near the typical VaR level, so the ranking is not uniform over all consistent quantile scores.

Table OA.9: Conditional predictive ability, dynamic quantile test, and loss-differential dispersion on the frontier rows.

<i>Giacomini–White CPA of d_{it} on $(1, \text{score}_{i,t-1})$</i>				
Comparison	Score	Wald χ^2_2	Slope of d on score	Slope t
GARCH- t minus engine	composite pct.	73.6	0.0081	8.6
GARCH- t minus engine	mk ₆₃ pct.	90.0	0.0068	9.8
Pooled GARCH-EVT minus engine	composite pct.	64.5	0.0072	7.5
GARCH- t minus pooled body	composite pct.	91.4	0.0105	9.8
<i>Dynamic quantile test, per-name pass rate at 5% (200 names)</i>				
Model	$\alpha = 1\%$		$\alpha = 2.5\%$	
Engine (accuracy layer)	86%		70%	
Engine + static shift (97.5%)	86%		84%	
GARCH(1,1)- t	84%		68%	
GARCH-EVT, pooled tail	83%		62%	
GARCH-EVT, per name	82%		60%	
GARCH + pooled body (no EVT)	69%		54%	
<i>Per-date loss differential, GARCH-t minus engine (pinball units)</i>				
Region	Win rate	Median	5th–95th pct.	SD
Top mk ₆₃ decile	74%	0.0102	-0.0258 to 0.0493	0.0265
Deciles 1–9	55%	0.0007	-0.0114 to 0.0096	0.0072
Overall	62%	0.0016	-0.0102 to 0.0114	0.0076

Same 200 names and 221,600 rows as Table OA.7. Win rate is the share of the 1,108 test dates on which the date-averaged engine loss is below the reference. The pooled dynamic quantile test rejects for every model at both levels, as pooled asset-day tests do at this sample size (Section 5.1 of the paper).

All benchmarks on the canonical rows

The GARCH-EVT comparison above is one row of a larger exercise (`job.bench_all.py`, result file `bench_all_results.json`): every model in the paper’s comparison set is scored on the identical 221,600 test rows of the 200-name panel, through every statistic. GARCH- t , GJR-GARCH-skew- t , the two GARCH-EVT forms, the GAS and Taylor models and SAV-CAViaR are estimated per name on the estimation window; the FHS variants take the empirical quantiles of the GARCH- t residuals (pooled across names, per name on the training window, or over a rolling 500-day window); EWMA is the RiskMetrics recursion with Gaussian quantiles; historical simulation is the rolling 500-day empirical return quantile. GAS and Taylor exist only at the two regulatory levels and enter the joint-score, dynamic-quantile and Murphy statistics. Table OA.10 gives the eleven-level pinball frontier, Table OA.11 the joint score, and Table OA.12 the conditional-predictive-ability, dynamic-quantile, Murphy and win-rate statistics.

Table OA.10: Every benchmark on the frontier rows: eleven-level pinball edge over GARCH- t by score region, engine head-to-head, and Model Confidence Set membership.

Model	Top mk ₆₃ decile	Deciles 1–9	Overall	Engine vs model (top)	MCS
Engine (accuracy layer)	+2.48 (10.08)	−0.01 (−0.25)	+0.27 (4.62)	—	
GARCH + pooled body (no EVT)	+2.98 (10.48)	+0.02 (0.43)	+0.35 (5.35)	−0.51 (−9.04)	✓
GARCH(1,1)- t	(reference)	(reference)	(reference)	+2.48 (10.08)	
GJR-GARCH-skew- t	+0.16 (0.60)	−0.01 (−0.13)	+0.01 (0.35)	+2.32 (9.70)	
GARCH-FHS (pooled)	−0.14 (−2.45)	−0.12 (−4.39)	−0.13 (−4.62)	+2.62 (8.86)	
GARCH-FHS (per name)	+0.34 (3.87)	−0.27 (−6.86)	−0.20 (−5.33)	+2.15 (9.67)	
GARCH-FHS (rolling 500d)	−0.75 (−6.68)	−0.03 (−0.55)	−0.11 (−2.19)	+3.21 (10.22)	
GARCH-EVT (pooled tail)	−0.16 (−2.69)	−0.12 (−4.55)	−0.12 (−4.89)	+2.64 (8.91)	
GARCH-EVT (per name)	+0.31 (3.62)	−0.22 (−6.61)	−0.16 (−5.00)	+2.18 (9.63)	
SAV-CAViaR	+2.81 (8.16)	+0.08 (1.60)	+0.39 (6.29)	−0.33 (−1.82)	✓
EWMA (RiskMetrics)	−7.31 (−9.91)	−0.43 (−3.54)	−1.20 (−8.39)	+9.13 (12.07)	
Historical simulation (500d)	+0.97 (1.88)	−1.98 (−8.39)	−1.65 (−7.04)	+1.53 (4.46)	

200 names, 221,600 test rows, the rows of Table 1 of the paper. Edge = (GARCH- t pinball − model pinball)/GARCH- t pinball over the eleven-level grid, per-date Newey–West(10) DM in parentheses. The last column is the 90% Model Confidence Set on the per-date eleven-level pinball series ($T_{\max}$, stationary bootstrap over dates, mean block 10, $B = 1,000$).

Expected-shortfall integral and an external transfer check

Two items recorded for completeness. (i) ES_{α} is the converged integral of the min-envelope curve (Section 2.2 of the paper). In one run on the same rows, replacing it by the GPD closed form of McNeil and Frey (2000) raises the accuracy layer’s 2.5% FZ0 from 1.8481 to 1.8508 (DM 4.2 in favour of the integral) and lowers its DM over GARCH- t from 5.5 to 4.0; at 1% the FZ0 rises from 2.1464 to 2.1487 (DM 2.5) and the DM over GARCH- t falls from 4.7 to 4.3. The 20-node midpoint rule of earlier versions understated the engine’s $|ES|$ by 1.65% at 1% and 1.66% at 2.5% on average, for every one of the 200 names, and the rolling-window empirical rows of Table 4 of the paper by 4–5%; those rows now use the exact tail mean. (ii) As a transfer check outside finance, the same

Table OA.11: Every benchmark on the full-panel FZ0 rows.

Model	$\alpha = 1\%$				$\alpha = 2.5\%$			
	FZ0	Breach	DM	MCS	FZ0	Breach	DM	MCS
Engine (accuracy layer)	2.1465	1.04%	—	✓	1.8486	2.78%	—	✓
Engine + static shift (97.5%)	2.1465	1.04%	0.00	✓	1.8664	1.66%	2.57	
GARCH + pooled body (no EVT)	2.1411	1.32%	-1.09	✓	1.8432	3.25%	-1.54	✓
GARCH(1,1)- t	2.1557	1.03%	4.94		1.8542	2.71%	5.50	
GJR-GARCH-skew- t	2.1505	1.02%	1.24	✓	1.8486	2.64%	-0.02	✓
GARCH-FHS (pooled)	2.1557	1.14%	3.93		1.8598	3.00%	5.89	
GARCH-FHS (per name)	2.1724	1.24%	4.91		1.8624	3.08%	5.35	
GARCH-FHS (rolling 500d)	2.1847	1.10%	6.61		1.8611	2.50%	4.10	
GARCH-EVT (pooled tail)	2.1551	1.06%	4.51		1.8583	2.89%	6.33	
GARCH-EVT (per name)	2.1693	1.13%	5.14		1.8596	2.93%	5.36	
GAS-FZ (PZC)	2.2324	0.92%	9.91		1.8990	2.53%	8.78	
ES-CAViaR (Taylor)	2.1647	0.98%	3.76		1.8490	2.62%	0.13	✓
EWMA (RiskMetrics)	2.3211	1.61%	6.89		1.9035	2.73%	6.25	
Historical simulation (500d)	2.2512	0.87%	6.34		1.9182	2.06%	5.28	

FZ0 is the zero-homogeneous Fissler–Ziegel joint (VaR, ES) loss, lower is better; DM > 0 means the row model scores worse than the accuracy layer (per-date Newey–West(10)). ES: closed forms for the t and Gaussian entrants; an exact substitution integral for the skew- t ; empirical tail means for HS and FHS; McNeil–Frey closed form for GARCH-EVT; the FZ0-estimated dynamic pair for GAS and Taylor (the GAS fit of record is this run’s); the converged integral of the monotonized curve for the engine (Section 2.2 of the paper), and a 20-node integral for the body, which has no quantile below its lowest fitted level. The static-shift row is the split-conformal shift of earlier versions, estimated on the body’s calibration errors; Stage 4 now uses the adaptive overlay of Section 5.1 of the paper. At 2.5% the engine, the body, GJR-GARCH-skew- t and the Taylor model form a tie group (FZ0 1.8432–1.8490, |DM| ≤ 1.54), reported as a group rather than an order. SAV-CAViaR produces no ES and is absent here. MCS: 90% Model Confidence Set on the per-date FZ0 series.

Table OA.12: Conditional predictive ability, dynamic quantile test, Murphy diagrams and win rates for every benchmark against the accuracy layer.

Model	CPA vs engine		DQ pass rate		Murphy: engine better		Engine win rate	
	Wald χ^2	slope t	1%	2.5%	1%	2.5%	top decile	overall
Engine + static shift (97.5%)	—	—	86%	84%	0%	45%	—	—
GARCH + pooled body (no EVT)	97.2	-10.1	69%	54%	82%	72%	12%	26%
GARCH(1,1)- t	73.6	8.6	84%	68%	70%	65%	74%	62%
GJR-GARCH-skew- t	87.7	9.5	90%	77%	48%	40%	72%	62%
GARCH-FHS (pooled)	63.4	7.2	78%	60%	88%	85%	73%	62%
GARCH-FHS (per name)	80.3	4.6	73%	55%	92%	90%	72%	62%
GARCH-FHS (rolling 500d)	138.2	12.0	84%	82%	85%	50%	74%	64%
GARCH-EVT (pooled tail)	64.5	7.5	83%	62%	85%	85%	74%	62%
GARCH-EVT (per name)	68.2	5.2	82%	60%	78%	72%	72%	62%
SAV-CAViaR	2.6	-0.2	78%	66%	80%	60%	52%	56%
GAS-FZ (PZC)	—	—	92%	82%	50%	78%	—	—
ES-CAViaR (Taylor)	—	—	78%	72%	70%	42%	—	—
EWMA (RiskMetrics)	172.6	13.8	29%	56%	95%	78%	83%	71%
Historical simulation (500d)	89.3	-0.4	48%	51%	48%	60%	58%	72%
Engine (accuracy layer)	—	—	86%	70%	—	—	—	—

CPA: Giacomini–White Wald test of equal conditional accuracy of the model and the engine with instruments (1, lagged composite score percentile), Newey–West(10) over per-date sums, and the date-clustered t of the slope of the loss differential (model minus engine) on the score; a positive slope means the engine’s advantage grows with the score. DQ: Engle–Manganelli dynamic quantile test with four hit lags and the VaR, share of names passing at 5%. Murphy: share of a forty-point return-space threshold grid on which the engine has the lower elementary quantile score. Win rate: share of test dates on which the date-averaged engine pinball is below the model’s. Models with an entry only at the regulatory levels (GAS, Taylor) have no pinball-based columns.

amortized pipeline scores a leakage-safe weighted scaled pinball of 0.269 (benchmark tier) on the M5 competition data.

OA.6 Frontier robustness: calendar overlap, family-wise error, and the universe rule

This section collects the constructions summarized in Section 3 of the paper.

Calendar overlap and the pooled learner. The design panels are split per name, 60/40 in each name’s own time index. In an unbalanced panel such a rule can place one name’s training window on another name’s test dates, so a pooled learner could meet a systemic event in training that a per-name model would meet only out of sample, the panel form of overlap leakage that purging and embargo address. Near-complete histories put the split points on almost the same date, so the channel is nearly degenerate here, and a strict calendar-time split closes it: each stage (each GARCH fit, the pooled learner, the GPD tail) is re-estimated strictly before 2020-01-01 and tested from that date for every name simultaneously. The frontier survives, with the top decile at +2.47% (per-date DM 6.22 over 1,258 dates; per-name split +2.7%, DM 6.5) and the overall edge

Table OA.13: The ten-day extension against iterated benchmarks (main text, ten-day subsection). Direct multi-day model’s pinball edge over each benchmark, $h = 10$, both eras, on the purged rows of the ten-day sections.

Benchmark	Design era, 2020–2024 test window		Holdout, mid-2008 to 2013	
	Edge % (DM; NO)	Cov. 1% / 2.5%	Edge % (DM; NO)	Cov. 1% / 2.5%
$\sqrt{h}$ -scaled GARCH- t (paper)	+1.41 (6.8; 4.3)	1.24 / 3.30	+0.52 (0.09; -0.02)	1.26 / 2.80
Iterated-variance GARCH- t	+0.79 (4.0; 2.7)	1.03 / 2.85	+0.54 (0.14; 0.15)	1.23 / 2.74
Simulated GARCH- t paths	+0.70 (3.5; 2.5)	1.20 / 2.93	+0.76 (0.25; 0.20)	1.44 / 2.79
Bootstrapped filtered paths	+0.51 (2.5; 1.1)	1.01 / 2.58	+0.81 (0.25; 0.31)	1.29 / 2.62
Envelope, min(direct, $\sqrt{h}$)	+1.14 (3.9; 2.4)	0.94 / 2.45	+0.15 (-0.32; 0.04)	1.01 / 2.38
Variance-ratio rescaling	-0.12 (-1.7; -0.3)	0.97 / 2.44	+0.07 (2.3; 2.0)	1.23 / 2.79
Direct model	—	1.13 / 2.80	—	1.24 / 2.82

Design era: 150 names, 164,850 test rows, 1,099 dates; holdout: 200 names, 278,808 rows, 1,486 dates. Edge = (benchmark pinball – direct pinball)/benchmark pinball over the eleven levels of the paper. Labels overlap (each row’s target is the next ten days), so the DM statistic is computed on per-date mean loss differences with a Newey–West lag of 10; NO repeats it on every tenth test date (110 and 149 dates). Coverage is the share of rows whose realized ten-day return falls below the forecast quantile. Iterated-variance GARCH- t uses the closed-form sum of the fitted GARCH variance forecasts over the ten days with the one-day t shape; the path benchmarks simulate 4,000 ten-day paths per row from the fitted GARCH- t with t innovations or with residuals bootstrapped from the training window. The ratio of the iterated ten-day variance to h times the one-day variance averages 1.07 in 2020, 1.11 in 2021, 1.06 in 2022, 1.12 in 2023, 1.11 in 2024 in the design era and 0.99 in 2008, 1.00 in 2009, 1.02 in 2010, 1.02 in 2011, 1.04 in 2012, 1.04 in 2013, which is why the $\sqrt{h}$ rule under-scales in the later era only. Holdout standard errors are roughly 3–6% of loss, so the out-of-era accuracy comparison is uninformative; see `tenday.diag.results.json`. The envelope (the level-wise minimum of the direct and $\sqrt{h}$ curves, rearranged) and the variance-ratio rescaling (the direct curve rescaled by the square root of the iterated-to- $\sqrt{h}$ variance ratio) come from a rerun on the same rows, `tenday.envelope.results.json`, in which the $\sqrt{h}$ row reads +1.39 (DM 6.7) in era and +0.51 (DM 0.05) out of era. In 2008–2009 the envelope breaches at 1.61% and 3.52%, against 1.93% and 3.99% for $\sqrt{h}$ scaling and 2.15% and 4.33% for the direct model. Both constructions were run with predictions written beforehand; the envelope did not meet its two, and the rescaling is exploratory.

Table OA.14: Learner feature ablation (Section 4 of the paper). The Table 1 learner re-trained with reduced state vectors; the score is computed as in the paper and is never a learner input.

Learner	Top mk ₆₃	Top sk ₆₃	Top composite	Deciles 1–8	Overall
Full (logsig, z_{t-1} , $ z _5$, sd ₂₁ , dn ₅)	+2.98 (10.5)	+2.78 (9.8)	+2.79 (8.9)	-0.20 to +0.11	+0.35 (5.3)
No dispersion (logsig, z_{t-1} , dn ₅)	+1.93 (10.4)	+1.78 (9.9)	+2.02 (10.2)	-0.40 to +0.08	+0.14 (2.3)
Scale only (logsig)	+1.63 (11.0)	+1.56 (10.6)	+1.81 (10.0)	-0.43 to +0.12	+0.12 (2.4)

200 CRSP names, 221,600 test asset-days, the rows of Table 1. Edge = (GARCH- t pinball – learner pinball)/GARCH- t pinball over the eleven levels, per-date DM with Newey–West lag 10 in parentheses. Against the full learner in the top mk₆₃ decile the no-dispersion learner loses 1.08% (DM -7.3) and the scale-only learner 1.39% (DM -8.2). Features: logsig is the log GARCH scale, z_{t-1} the lagged standardized residual, $|z|_5$ its five-day mean absolute value, sd₂₁ its 21-day standard deviation, dn₅ the five-day share of down days. Result file `feature.ablation.results.json`; predictions in the script header.

Table OA.15: A state-adaptive parametric benchmark (Section 4 of the paper): GARCH- t with a time-varying degrees-of-freedom law, on the rows of Table 1.

Comparison	Top mk ₆₃ decile	Top composite decile	Deciles 1–9	Overall
Adaptive- ν GARCH- t over GARCH- t	+0.27 (1.9)	+0.03 (0.8)	+0.07 (3.8)	+0.10 (3.8)
Learner over GARCH- t (Table 1)	+2.98 (10.5)	+2.79 (8.9)	+0.02 (0.4)	+0.35 (5.3)
Learner over adaptive- ν GARCH- t	+2.72 (9.7)	+2.76 (8.1)	−0.05 (−0.7)	+0.26 (3.4)

The benchmark keeps the paper’s GARCH(1,1)- t scale and replaces the constant ν by $\nu_t = 2.1 + \exp(h_t)$, $h_t = c + a|z_{t-1}| + b h_{t-1}$, an autoregressive conditional density in the spirit of Hansen (1994), fitted by maximum likelihood on each name’s training residuals with the scale held fixed; forecast quantile $\hat{\mu} + \hat{\sigma}_t t_\tau(\nu_t)/\sqrt{\nu_t/(\nu_t - 2)}$. Median fitted $a = 0.11$, $b = 0.64$; median training log-likelihood gain over constant ν is 2.9 (48% of names gain more than 3); the typical name’s ν_t runs from 3.7 to 5.2 between its 10th and 90th percentiles. Edge = (first model pinball − second model pinball)/first, eleven levels, per-date DM with Newey–West lag 10. Result file `arcd_bench_results.json`; predictions in the script header.

Table OA.16: Synthetic truth (Section 4 of the paper): the estimator against the parametric model on its own scale, on panels simulated from a known model.

Data generated by	Scale of both models	Pinball DM		FZ0 difference ($\times 10^3$)	
		Overall	Top mk ₆₃ decile	1%	2.5%
GARCH- t	GARCH- t (true)	−5.20	−1.85	+0.27	+0.01
	Beta- t -EGARCH	10.17	7.71	−6.41	−5.37
Beta- t -EGARCH	Beta- t -EGARCH (true)	−4.29	−1.41	+0.29	−0.18
	GARCH- t	10.32	8.14	−1.62	−1.61

Each panel is simulated per name from the named model, 200 names and 221,600 test rows. In each row the estimator (pooled shape learner and EVT tail on the stated scale) is compared with the parametric Student- t model on the same scale. Pinball DM is the per-date Diebold–Mariano statistic of the estimator over that model on the eleven-level grid (positive favors the estimator); the FZ0 columns give mean FZ0 of the estimator minus that of the parametric model (negative favors the estimator), with the estimator’s ES the converged integral and the parametric ES in closed form. Against the true model the estimator is slightly worse on pinball (by 0.02% in both designs) and within 3×10^{-4} on FZ0; against the misspecified scale’s parametric tail it wins by $1.6\text{--}6.4 \times 10^{-3}$. The ES convention matters here: these gaps are smaller than the error of a 20-node midpoint ES, and under a correctly specified model the McNeil–Frey test with a 20-node ES rejects in 22.5% of simulated samples at a nominal 5% (6% with the exact ES). Result file `synthetic_truth_results.json`.

+0.71% (DM 1.87). The bottom decile is elevated at +1.3% because a 2020–22 test era leaves few calm cells, but the top decile roughly doubles each other cell. Inside the 2000–2013 holdout, re-estimating strictly before 2007-01-01 and testing on 2007–2013 removes the 2008 crisis from every training window; the pooled learner still wins overall (+0.29%, per-date DM 3.08 over 1,762 dates) and the decile profile keeps its shape (top three deciles +0.45–+1.04% against +0.06–+0.09%). The annual walk-forward refits the per-name GARCH on expanding data at each January-1 cutoff from 2020 to 2024, rebuilds the residuals and score, and retrains the learner; the top decile carries +2.12% at DM 4.41 while the overall edge is no longer detectable (−0.22%, DM −0.59).

Family-wise error over the signal-decile grid. The frontier grid spans three signals and ten deciles. Romano–Wolf step-down family-wise control over the full 3×10 grid takes studentized per-date mean edges per cell, a stationary bootstrap over dates that draws the same dates for every cell (mean block 10 days, $B = 2,000$, block 20 as sensitivity), and one-sided step-down adjusted p -values. Nine of the thirty cells survive under both block lengths, all frontier cells: the top deciles of the three signals (mk_{63} $t = 9.4$, skew_{63} $t = 9.2$, jump_5 $t = 5.1$, adjusted $p < 0.001$), the ninth deciles of mk_{63} (adjusted $p = 0.022$) and skew_{63} , skew_{63} ’s seventh, and jump_5 ’s U-shaped far cells. Mid-grid cells with an unadjusted t near two lose significance under the adjustment.

The point-in-time universe. The holdout’s observation filter conditions on names present for most of the era, tilting the universe toward survivors. A point-in-time re-selection fixes the universe at the top 200 names by market capitalization as of Q1-2000, applies no minimum-history filter, and rides each delisting return to the end. On the 164 of 200 names that clear the symmetric fitting floors (36 dropped, 79 delisting returns ridden to the end), the overall edge is +0.58% at per-date DM 8.96 over 2,692 dates, the frozen design-era top bucket carries +2.94% at DM 6.6 (5.9% occupancy), and the bottom bucket is +0.44% (DM 5.4). The edge is larger than in the survivor-tilted holdout, so the observation filter biased the reported edge downward.

OA.7 Estimator schematic and evaluation panels

Table OA.17: Sample flow across evaluation panels (Section 3 of the paper).

Panel	Source	Era	Names	Test obs	Role
Design equity panel	CRSP/WRDS	2014–2024	200	221,600	frontier map, estimator development
Untouched-era holdout	CRSP/WRDS	2000–2013	200	280,608	frozen-spec replication
Subpanel	CRSP/WRDS	2014–2024	140	$\approx 155,000$	neural-network comparison, auxiliary checks
Cross-asset sweep	Bloomberg	through 2026	43		^a frontier breadth
Country indices	Bloomberg	through 2026	26		^a frontier breadth

^a Frozen exhibits. Per-instrument windows vary and are documented instrument-by-instrument in the repository.




Figure OA.8: The residual-hybrid estimator as a pipeline (Section 2 of the paper). Solid arrows carry data; the dashed top arrow carries the conditional scale $\hat{\sigma}_t$ that rescales the finished residual quantile curve. Dashed boxes are inputs shared across assets (the pooled panel that trains the amortized shape model) or observed as the forecasts are scored (the breaches behind the adaptive overlay).

OA.8 The two-regime example

A synthetic boundary check: when the frontier does *not* appear

Before the real data, a simulation with known truth (run once, seed fixed) checks the formal results and, more usefully, delineates when the frontier edge should *fail* to appear. The data-generating process is deliberately favorable to the score: unit-variance Student- t residuals whose degrees of freedom switch between a calm regime ($\nu = 12$) and a fat regime ($\nu = 3$) under a persistent Markov chain; forecasters see the trailing 63-day kurtosis score and compete at $\tau \in \{0.01, 0.025\}$ (80,000 days; scale known, isolating shape). First, the Lemma 1 of the paper identity verifies numerically: the regret integral and the Monte-Carlo regret agree to five decimals (3.5×10^{-5} at $\tau = 0.01$). Second, the conformal shift moves realized coverage toward nominal on the held-out split, as Proposition 1 of the paper promises. Third, and most instructive, the score-binned nonparametric forecaster does *not* beat the fixed shape here: a single $t_{\hat{\nu}}$ fitted on the mixture ($\hat{\nu} = 6.8$) is nearly unbeatable at these levels, with total regret under 0.5% and no decile gradient. Two mechanisms explain it. (i) By Lemma 1 of the paper regret is *quadratic* in the quantile wedge, and at $\tau \geq 0.01$ the wedge between the fitted compromise and either regime is small, so the regret from one fitted shape across both regimes is second order. (ii) The sample kurtosis of a t_ν law is a noisy statistic when tails are fat: its usual moment-based asymptotic variance requires a finite eighth moment ($\nu > 8$), and the population kurtosis itself does not exist for $\nu \leq 4$, so sixty-three observations cannot reliably rank regimes that differ only in a static ν . The jump and asymmetry signals, whose top deciles independently survive familywise control, corroborate the ordering, and a tail-index (Hill-type) variant of the score is an obvious robustness extension. For the empirical sections this means the real-data top-decile edge of Section 4 of the paper is not what averaging

stationary fat tails produces in this simulation (a fixed shape handles that configuration) and is consistent with *local* shape breaks a single fitted ν cannot straddle: asymmetry, jump clustering, and residual misfit that co-move with the score. This is the same conclusion the Korea 2026 falsification case reaches from the opposite direction (price crashes are not residual misspecification), and it is why the score combines kurtosis with asymmetry and a jump indicator rather than kurtosis alone. (Simulation script `toy_example.py` ships with the code release.)

OA.9 Trading implications

Trading implications. The universes where the nonparametric tail wins (crisis FX, volatility indices, credit spreads, frontier equities) are the desks where tail-shape-aware sizing and hedging matter most. The *direction* of the trade matters as much as its location: a correct physical-measure tail forecast does not by itself make bought crash insurance profitable, because the crash risk premium can survive conditioning on a correct forecast. A possible application is the opposite side, selling tail protection with the model supplying tail control rather than alpha on the mean. The paper evaluates no prices, premia, costs or trading outcomes, so the application is untested here. The gating experiment provides no evidence that the score times individual crash days, so its role there would be continuous tail-risk measurement.

OA.10 Proofs

The three formal results of Sections 2.2 and 2.4 of the paper are stated in the body; the complete arguments, with all regularity hypotheses made explicit, are collected here. Throughout, $\rho_\tau(u) = u(\tau - \mathbb{I}\{u < 0\})$ and $q_F = F^{-1}(\tau) = \inf\{x : F(x) \geq \tau\}$ is the lower generalized inverse.

Proof of Lemma 1 of the paper (pinball regret identity)

Assume $\mathbb{E}|Z| < \infty$, so

$$S(q) = (1 - \tau) \int_{(-\infty, q]} (q - z) dF(z) + \tau \int_{(q, \infty)} (z - q) dF(z)$$

is finite for every q , and S is convex (an expectation of the convex maps $q \mapsto \rho_\tau(z - q)$). For fixed z , $q \mapsto \rho_\tau(z - q)$ is Lipschitz with constant $\max(\tau, 1 - \tau)$, so the difference quotients $[\rho_\tau(Z - q') - \rho_\tau(Z - q)]/(q' - q)$ are bounded by the integrable constant $\max(\tau, 1 - \tau)$; dominated convergence justifies differentiating under the expectation wherever the one-sided limits exist. Off the atoms of F ,

$$\frac{d}{dq} \rho_\tau(z - q) = \mathbb{I}\{z < q\} - \tau, \quad \implies \quad S'(q) = \mathbb{E} \mathbb{I}\{Z < q\} - \tau = F(q^-) - \tau,$$

which equals $F(q) - \tau$ for all but countably many q . Since S is convex, hence absolutely continuous, the fundamental theorem of calculus gives

$$S(q) - S(q_F) = \int_{q_F}^q (F(u) - \tau) du,$$

the countable atom set being Lebesgue-null; *no continuity hypothesis on F is needed for this identity*. For nonnegativity, note $q_F = \inf\{x : F(x) \geq \tau\}$ implies $F(u) \geq \tau$ for $u \geq q_F$ and $F(u) < \tau$ for $u < q_F$. If $q > q_F$ the integrand is ≥ 0 on $[q_F, q]$, so the integral is ≥ 0 ; if $q < q_F$,

$$S(q) - S(q_F) = \int_{q_F}^q (F(u) - \tau) du = \int_q^{q_F} (\tau - F(u)) du \geq 0,$$

the integrand again ≥ 0 on $[q, q_F]$. Equality holds iff $F(u) = \tau$ for Lebesgue-a.e. u between q_F and q . Finally, if F is absolutely continuous with density $0 < m \leq f \leq M$ on the interval joining q_F and q , then F is continuous there, so $F(q_F) = \tau$ and $F(u) - \tau = \int_{q_F}^u f$; hence $m|u - q_F| \leq |F(u) - \tau| \leq M|u - q_F|$, and integrating over the interval yields $\frac{m}{2}(q - q_F)^2 \leq S(q) - S(q_F) \leq \frac{M}{2}(q - q_F)^2$. $\square$

Proof of Proposition 1 of the paper (split-conformal validity)

Let $k = \lceil (n+1)\tau \rceil$ and assume $1 \leq k \leq n$ so that the k -th order statistic $e_{(k)} = c_\tau$ of the n calibration errors $\{e_1, \dots, e_n\}$ exists; when $k = n+1$ (i.e. τ within $1/(n+1)$ of 1) set $c_\tau = +\infty$, giving coverage 1. Suppose the $n+1$ errors $e_1, \dots, e_{n+1}$ (calibration plus the test error $e_{n+1} = z_{n+1} - \widehat{q}(\tau \mid s_{n+1})$)

are exchangeable and almost surely distinct. Then the rank R of e_{n+1} among the pooled $n+1$ values is uniformly distributed on $\{1, \dots, n+1\}$. By construction $\tilde{q}(\tau \mid s_{n+1}) = \hat{q}(\tau \mid s_{n+1}) + c_\tau$, so

$$\{z_{n+1} \leq \tilde{q}(\tau \mid s_{n+1})\} = \{e_{n+1} \leq c_\tau\} = \{e_{n+1} \leq e_{(k)}\} = \{R \leq k\},$$

the middle equality using a.s. distinctness (with ties the middle equality fails; e.g. if all e_i are equal the event has probability 1, which can exceed the stated upper bound). Therefore

$$\Pr(z_{n+1} \leq \tilde{q}(\tau \mid s_{n+1})) = \Pr(R \leq k) = \frac{k}{n+1} = \frac{\lceil (n+1)\tau \rceil}{n+1}.$$

Since $(n+1)\tau \leq \lceil (n+1)\tau \rceil < (n+1)\tau + 1$, dividing by $n+1$ gives $\tau \leq k/(n+1) < \tau + \frac{1}{n+1}$. The probability is over the joint law of calibration and test points; the guarantee is marginal, not conditional on s_{n+1} , because a single scalar shift c_τ pooled across states cannot deliver state-conditional coverage. $\square$

Proof of Proposition 2 of the paper (tail wedge)

Fix $\nu > 2$, so the t_ν law has finite variance and its unit-variance version $Z = T_\nu \sqrt{(\nu-2)/\nu}$ is well defined. The t_ν density is $\sim C_\nu |x|^{-(\nu+1)}$ as $|x| \rightarrow \infty$, so the lower tail is regularly varying with index $-\nu$: $\Pr(T_\nu < -x) = \tilde{c}_\nu x^{-\nu}(1 + o(1))$. Unit-variance scaling multiplies the constant, $\Pr(Z < -x) = c_\nu x^{-\nu}(1 + o(1))$ with $c_\nu = \tilde{c}_\nu((\nu-2)/\nu)^{\nu/2}$, and leaves the index ν unchanged. Writing $U(\tau) = |Q_\nu(\tau)|$ and inverting the regularly varying tail (Embrechts et al., 1997) gives $\tau = c_\nu U(\tau)^{-\nu}(1 + o(1))$, hence $U(\tau) = (c_\nu/\tau)^{1/\nu}(1 + o(1))$ as $\tau \downarrow 0$, i.e. $Q_\nu(\tau) = -(c_\nu/\tau)^{1/\nu}(1 + o(1))$. For $\nu_1 < \nu_2$,

$$\frac{|Q_{\nu_1}(\tau)|}{|Q_{\nu_2}(\tau)|} = \frac{(c_{\nu_1}/\tau)^{1/\nu_1}}{(c_{\nu_2}/\tau)^{1/\nu_2}} (1 + o(1)) = \text{const} \cdot \tau^{-(1/\nu_1 - 1/\nu_2)} (1 + o(1)) \rightarrow \infty$$

because $1/\nu_1 > 1/\nu_2$. Thus for all sufficiently small τ both quantiles are negative with $|Q_{\nu_1}(\tau)| > |Q_{\nu_2}(\tau)|$, i.e. $Q_{\nu_1}(\tau) < Q_{\nu_2}(\tau)$, and the wedge $|Q_{\nu_1}(\tau) - Q_{\nu_2}(\tau)| \rightarrow \infty$. The ordering therefore holds for every $\tau < \delta$ for some $\delta > 0$; putting $\tau^* = \min(\delta, \frac{1}{2})$ gives $\tau^* \in (0, \frac{1}{2})$ as asserted. (This establishes *existence* of τ^* ; the actual crossover τ_c where the ordering flips is not claimed here and is located numerically in the crossover remark of the paper.) Under local truth t_{ν_1} , applying the corresponding equation of the paper with $q = Q_{\nu_2}(\tau)$ and $q_F = Q_{\nu_1}(\tau)$ bounds the regret below by $\frac{m}{2}(Q_{\nu_1}(\tau) - Q_{\nu_2}(\tau))^2$, with m the minimum of the t_{ν_1} density on the wedge interval. $\square$

Remark on the score ordering

We state the score-ordering as an empirical regularity rather than a theorem, and record here why no clean monotonicity theorem holds at the deployable levels. The tail-wedge result (Proposition 2

of the paper) orders a given pair of Student- t tail indices only below a pair-specific crossover level $\tau^*(\nu_1, \nu_2)$; it does not deliver monotonicity of the wedge in the tail index at a fixed operational τ . More fundamentally, the innovations are unit-variance by construction, and for the unit-variance t_ν the level- τ quantile is $Q_\nu(\tau) = \sqrt{(\nu - 2)/\nu} t_\nu^{-1}(\tau)$, which is *not* monotone in ν : as $\nu \downarrow 2$ the normalization $\sqrt{(\nu - 2)/\nu} \rightarrow 0$ pulls the fixed- τ quantile back toward zero, so an arbitrarily heavy-tailed unit-variance t eventually has a fixed- τ quantile arbitrarily close to the center rather than far into the tail. The non-monotonicity is already active at the regulatory levels: at $\tau = 0.025$, $Q_5(0.025) \approx -1.99$ is *less* extreme than $Q_8(0.025) \approx -2.00$, so a heavier-tailed standardized t carries the less extreme 2.5% quantile. A clean ordering theorem would therefore have to restrict the tail index to an explicit interval, isolate a range of τ on which $\partial Q_\nu(\tau)/\partial \nu$ has one sign, and control the density factor in the regret expansion; we do not pursue that narrow statement. The exact content we retain is the regret identity (Lemma 1 of the paper) and the tail-wedge ordering (Proposition 2 of the paper), and the frontier itself is documented empirically in the paper. $\square$

References

- Carlo Acerbi and Balázs Székely. Back-testing expected shortfall. *Risk*, 27(11):76–81, 2014.
- Giovanni Barone-Adesi, Kostas Giannopoulos, and Les Vosper. VaR without correlations for portfolios of derivative securities. *Journal of Futures Markets*, 19(5):583–602, 1999.
- Basel Committee on Banking Supervision. Minimum capital requirements for market risk. Technical report, Bank for International Settlements, 2019.
- Sebastian Bayer and Timo Dimitriadis. Regression-based expected shortfall backtesting. *Journal of Financial Econometrics*, 20(3):437–471, 2022.
- Tim Bollerslev. Generalized autoregressive conditional heteroskedasticity. *Journal of Econometrics*, 31(3):307–327, 1986.
- Peter F. Christoffersen. Evaluating interval forecasts. *International Economic Review*, 39(4):841–862, 1998.
- Will Dabney, Georg Ostrovski, David Silver, and Rémi Munos. Implicit quantile networks for distributional reinforcement learning. In *Proceedings of the 35th International Conference on Machine Learning (ICML)*, pages 1096–1105, 2018.
- Francis X. Diebold and Roberto S. Mariano. Comparing predictive accuracy. *Journal of Business & Economic Statistics*, 13(3):253–263, 1995.
- Werner Ehm, Tilmann Gneiting, Alexander Jordan, and Fabian Krüger. Of quantiles and expectiles: Consistent scoring functions, Choquet representations and forecast rankings. *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, 78(3):505–562, 2016.
- Paul Embrechts, Claudia Klüppelberg, and Thomas Mikosch. *Modelling Extremal Events for Insurance and Finance*. Springer, 1997.

- Robert F. Engle and Simone Manganelli. CAViaR: Conditional autoregressive value at risk by regression quantiles. *Journal of Business & Economic Statistics*, 22(4):367–381, 2004.
- Tobias Fissler and Johanna F. Ziegel. Higher order elicibility and Osband’s principle. *Annals of Statistics*, 44(4):1680–1707, 2016.
- Raffaella Giacomini and Halbert White. Tests of conditional predictive ability. *Econometrica*, 74(6):1545–1578, 2006.
- Lawrence R. Glosten, Ravi Jagannathan, and David E. Runkle. On the relation between the expected value and the volatility of the nominal excess return on stocks. *Journal of Finance*, 48(5):1779–1801, 1993.
- Bruce E. Hansen. Autoregressive conditional density estimation. *International Economic Review*, 35(3):705–730, 1994.
- Peter R. Hansen, Asger Lunde, and James M. Nason. The model confidence set. *Econometrica*, 79(2):453–497, 2011.
- Peter Reinhard Hansen. A test for superior predictive ability. *Journal of Business & Economic Statistics*, 23(4):365–380, 2005.
- Peter Reinhard Hansen, Zhuo Huang, and Howard Howan Shek. Realized GARCH: A joint model for returns and realized measures of volatility. *Journal of Applied Econometrics*, 27(6):877–906, 2012.
- J.P. Morgan/Reuters. RiskMetrics — technical document. Technical report, J.P. Morgan/Reuters, New York, 1996.
- Keith Kuuster, Stefan Mittnik, and Marc S. Paoletta. Value-at-risk prediction: A comparison of alternative strategies. *Journal of Financial Econometrics*, 4(1):53–89, 2006.
- Paul H. Kupiec. Techniques for verifying the accuracy of risk measurement models. *Journal of Derivatives*, 3(2):73–84, 1995.
- Alexander J. McNeil and Rüdiger Frey. Estimation of tail-related risk measures for heteroscedastic financial time series: An extreme value approach. *Journal of Empirical Finance*, 7(3–4):271–300, 2000.
- Natalia Nolde and Johanna F. Ziegel. Elicibility and backtesting: Perspectives for banking regulation. *Annals of Applied Statistics*, 11(4):1833–1874, 2017.
- Andrew J. Patton, Johanna F. Ziegel, and Rui Chen. Dynamic semiparametric models for expected shortfall (and Value-at-Risk). *Journal of Econometrics*, 211(2):388–413, 2019.
- Nicholas G. Polson and Vadim Sokolov. Generative AI for Bayesian computation. *arXiv preprint arXiv:2305.14972*, 2023.